



Recommender systems with Python & TensorFlow



TensorFlow



TensorFlow based Matrix Factorisation




$$\begin{bmatrix} \square & \square \\ \square & \square \\ \square & \square \\ \square & \square \end{bmatrix} \times \begin{bmatrix} \square & \square & \square & \square & \square & \square \\ \square & \square & \square & \square & \square & \square \end{bmatrix} \approx \begin{bmatrix} \square & \square & \square & \square & \square & \square \\ \square & \square & \square & \square & \square & \square \\ \square & \square & \square & \square & \square & \square \\ \square & \square & \square & \square & \square & \square \end{bmatrix}$$

Brief introduction to Recommender systems

From the Netflix prize (2008)

To real world use cases (Amazon, Spotify, Youtube, Criteo, Mendeley, Schibsted)



(From the [Mendeley's blog](#))

Collaborative Filtering for Implicit Feedbacks (2008)

- + Not only explicit ratings, also negative implicit interactions

$$\min_{x_*, y_*} \sum_{u, i} c_{ui} (p_{ui} - x_u^T y_i)^2 + \lambda \left(\sum_u \|x_u\|^2 + \sum_i \|y_i\|^2 \right)$$

- + Alternating Least Squares is fast and distributed (org.apache.spark.ml)

$$x_u = (Y^T C^u Y + \lambda I)^{-1} Y^T C^u p(u)$$

- RMSE only
- Click only, no user or item features => cold start issue!

Logistic MF for implicit feedback data (2014)

- + Binary click prediction

$$p(l_{ui} \mid x_u, y_i, \beta_i, \beta_j) = \frac{\exp(x_i y_i^T + \beta_u + \beta_i)}{1 + \exp(x_u y_i^T + \beta_u + \beta_i)}$$

```
class LogisticMF():  
    ...  
  
    def inference(user_ids, item_ids):  
  
        logits = np.dot(self.user_vectors[user_ids], self.item_vectors[item_ids].T) \  
            + self.user_biases[user_ids] + self.item_biases[item_ids].T  
  
        return 1 / (1 + np.exp(-logits))
```

Deriving gradients manually...

$$\frac{\partial}{\partial x_u} = \sum_i \alpha r_{ui} y_i - \frac{y_i (1 + \alpha r_{ui}) \exp(x_u y_i^T + \beta_u + \beta_i)}{1 + \exp(x_u y_i^T + \beta_u + \beta_i)} - \lambda x_u$$

```
def user_gradients():  
    vec_deriv = np.dot(self.counts, self.item_vectors)  
    bias_deriv = np.expand_dims(np.sum(self.counts, axis=1), 1)  
  
    A = np.exp(  
        np.dot(self.user_vectors, self.item_vectors.T) \  
        + self.user_biases + self.item_biases.T)  
    A /= (A + self.ones)  
    A = (self.counts + self.ones) * A  
  
    vec_deriv -= np.dot(A, self.item_vectors) - self.reg_param * self.user_vectors  
    bias_deriv -= np.expand_dims(np.sum(A, axis=1), 1)  
  
    return vec_deriv, bias_deriv
```

And the gradient descent step (ADAGRAD)

$$x_u^t = x_u^{t-1} + \frac{\gamma g_u^{t-1}}{\sqrt{\sum_{t'=1}^{t-1} g_u^{t',2}}}$$

```
def ada_grad_step():  
    user_vec_deriv, user_bias_deriv = self.user_deriv()  
    user_vec_deriv_sum += np.square(user_vec_deriv)  
    user_bias_deriv_sum += np.square(user_bias_deriv)  
    vec_step_size = self.gamma / np.sqrt(user_vec_deriv_sum)  
    bias_step_size = self.gamma / np.sqrt(user_bias_deriv_sum)  
    self.user_vectors += vec_step_size * user_vec_deriv  
    self.user_biases += bias_step_size * user_bias_deriv
```


Welcome to symbolic computing and auto-diff!



Declaring the inference step in TF

```
def user_bias(self, user_ids):
    with tf.name_scope('B_user'):
        return tf.squeeze(tf.nn.embedding_lookup(params=self.user_biases, ids=user_ids), name='B_user')

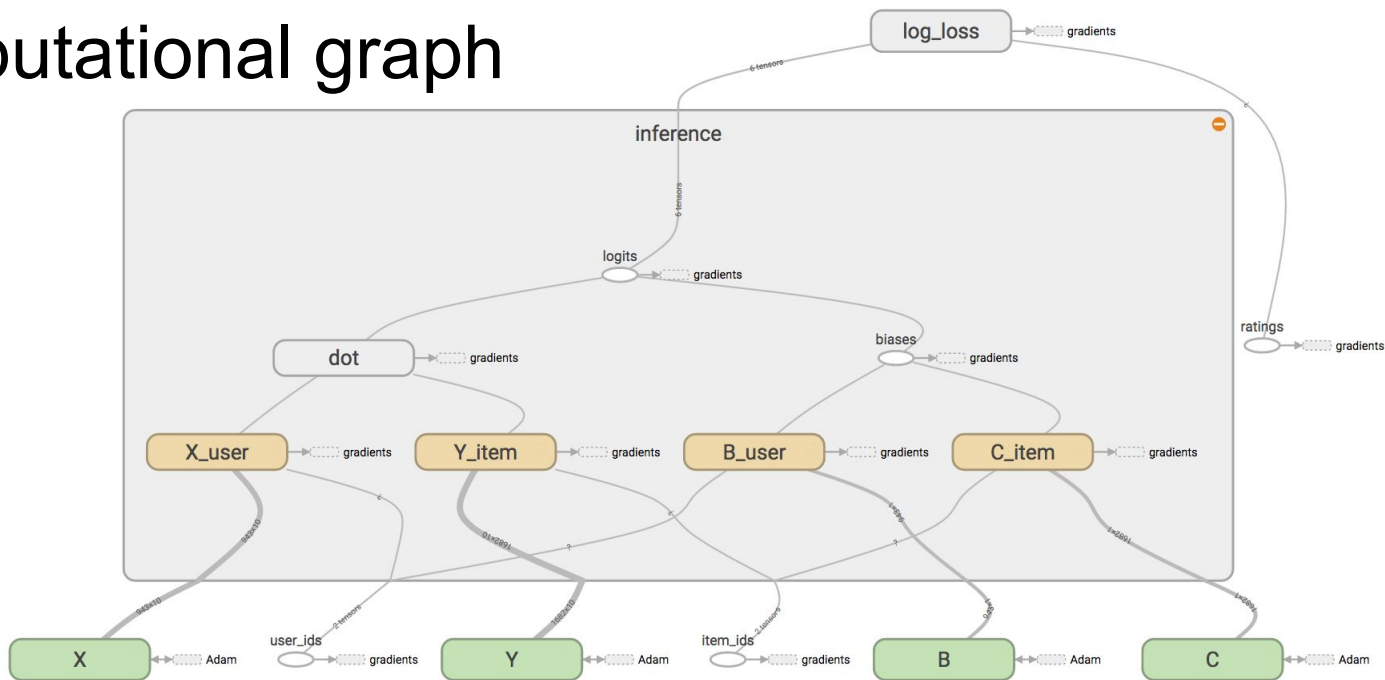
def item_bias(self, item_ids):
    with tf.name_scope('C_item'):
        return tf.squeeze(tf.nn.embedding_lookup(params=self.item_biases, ids=item_ids), name='C_item')

def user_item_product(self, user_ids, item_ids):
    with tf.name_scope('X_user'):
        batch_user_factors = tf.squeeze(tf.nn.embedding_lookup(self.user_factors, user_ids))
    with tf.name_scope('Y_item'):
        batch_item_factors = tf.squeeze(tf.nn.embedding_lookup(self.item_factors, item_ids))
    with tf.name_scope('dot'):
        factors_prediction = tf.reduce_mean(
            tf.mul(batch_user_factors, batch_item_factors), reduction_indices=1)
    return factors_prediction

def inference(self, user_ids, item_ids):
    with tf.name_scope('inference'):
        return tf.add(
            self.user_item_product(user_ids, item_ids),
            tf.add(self.user_bias(user_ids), self.item_bias(item_ids), name='biases'),
            name='logits')
```

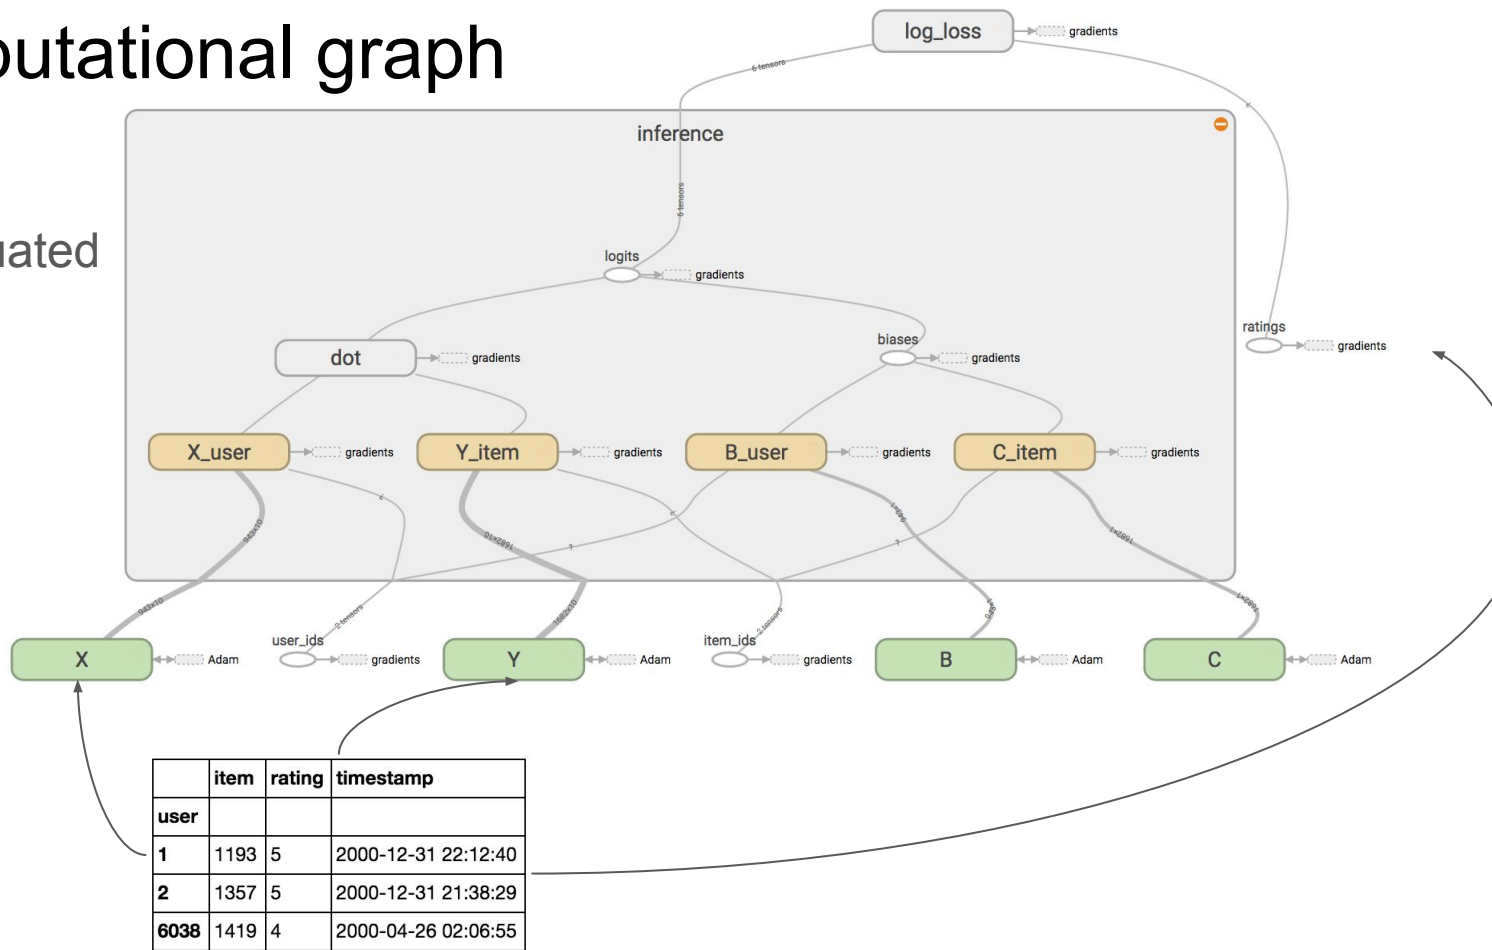
As a computational graph

- Compiled



As a computational graph

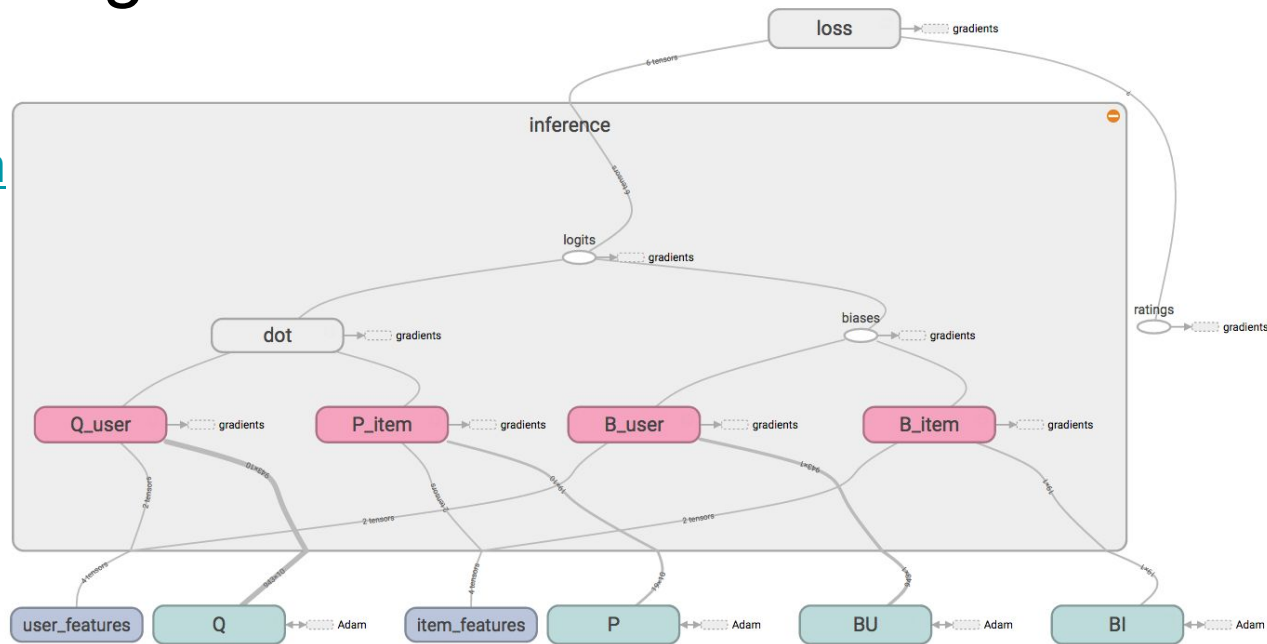
- Compiled
- Lazily evaluated by batches



Metadata Embeddings for user and item cold start

- + Incorporate metadata
- + github.com/lyst/lightfm

	item	rating	timestamp
user			
1	1193	5	2000-12-31 22:12:40
2	1357	5	2000-12-31 21:38:29
6038	1419	4	2000-04-26 02:06:55

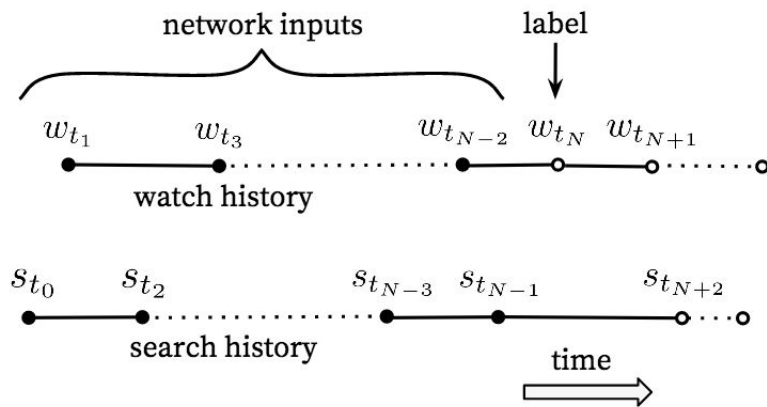


	age=1	age=18	age=25	age=35	age=45	age=50	age=56	gender=F	gender=M
user									
1	1	0	0	0	0	0	0	1	0
2	0	0	0	0	0	0	1	0	1
6038	0	0	0	0	0	0	1	1	0

$$q_u = \sum_{j \in f_u} e_j^U \quad p_i = \sum_{j \in f_i} e_j^I$$

Deep Neural Networks for YouTube Recs

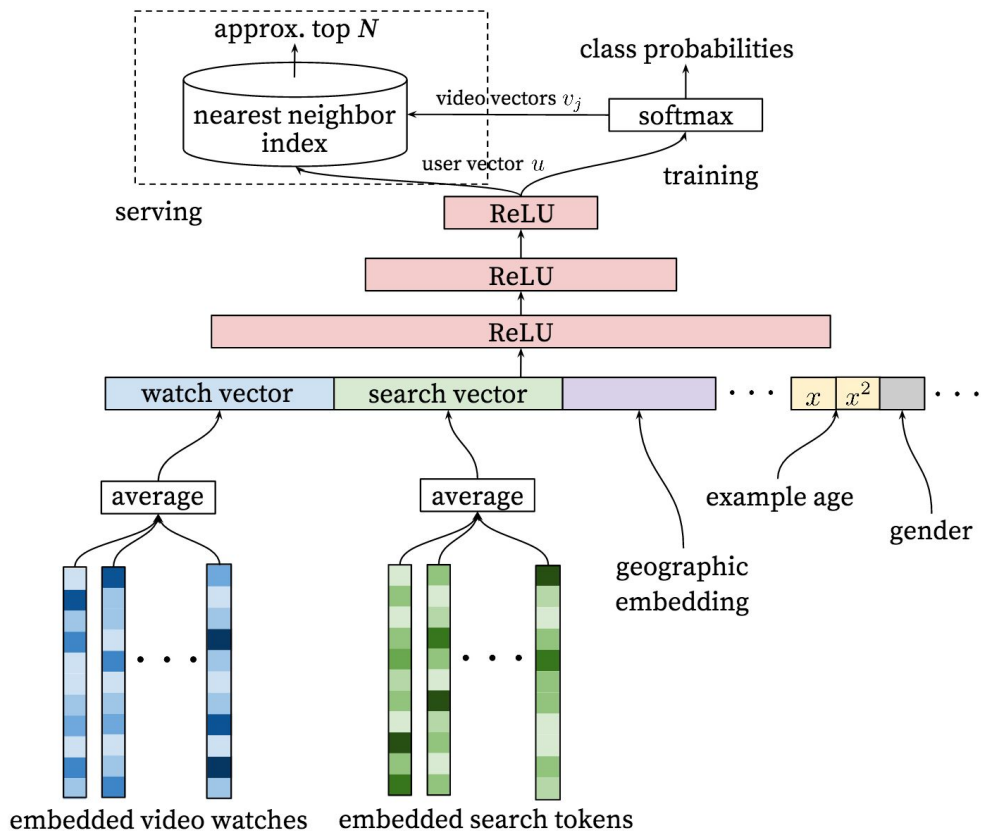
- + Predict future interactions from past interactions
- + Extreme multi-class classification “à la word2vec”



(b) Predicting future watch

$$P(w_t = i | U, C) = \frac{e^{v_i u}}{\sum_{j \in V} e^{v_j u}}$$

Deep Neural Networks for YouTube Recs

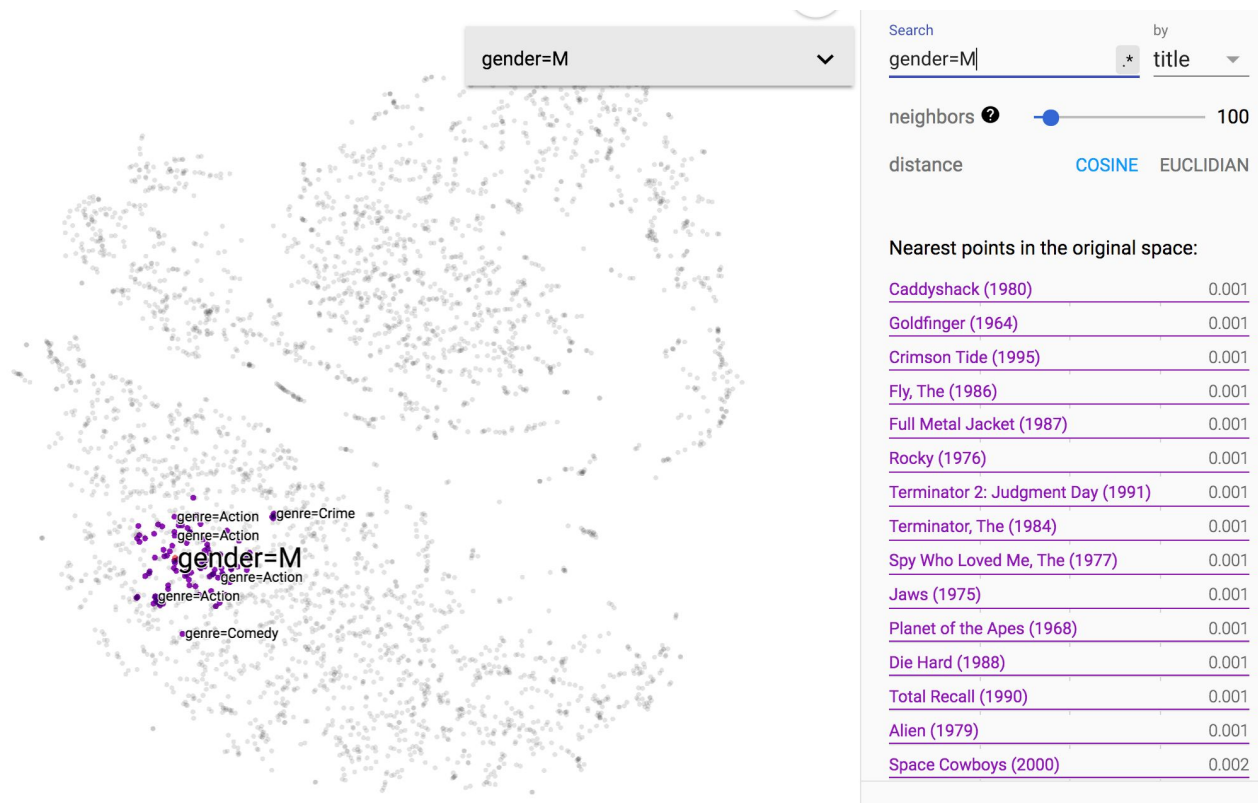


Embeddings visualisation with Tensorboard

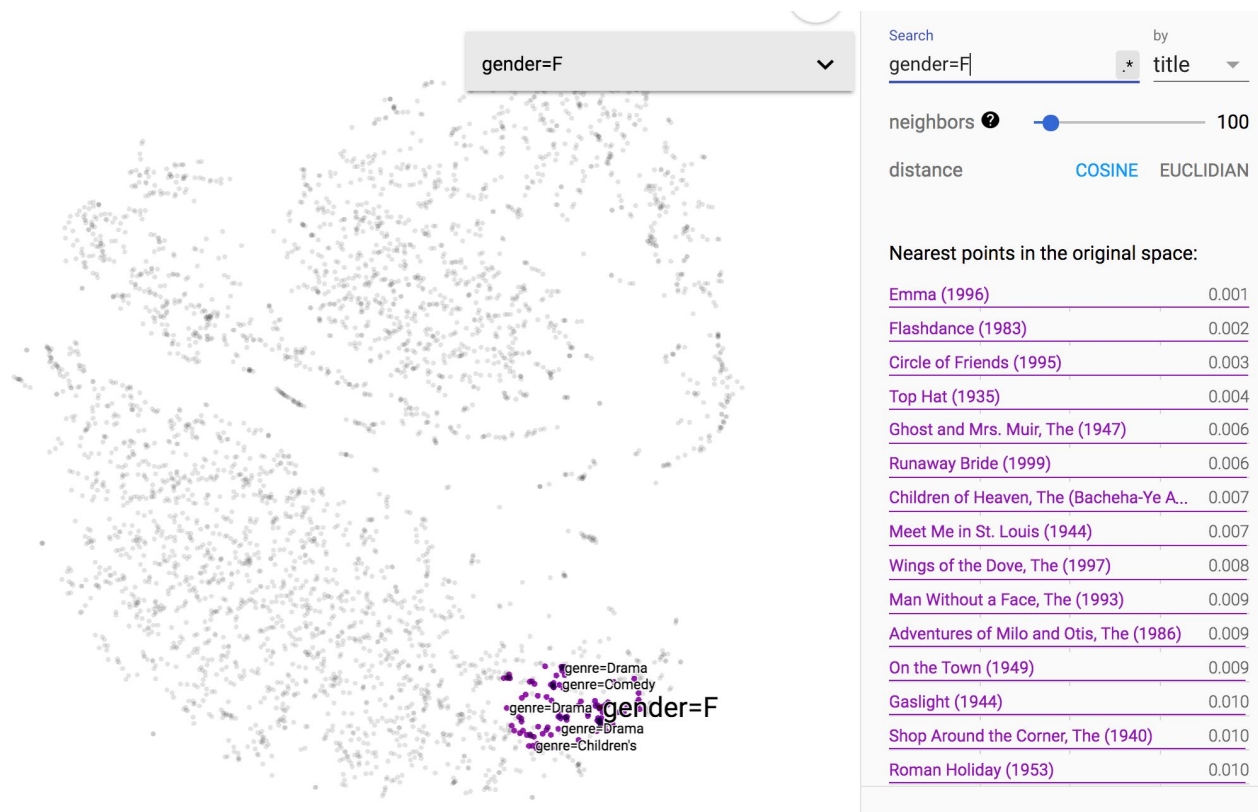


genre=Animation	90
genre=Adventure	155
genre=Comedy	1024
genre=Action	503
genre=Drama	1176
genre=Thriller	101
genre=Crime	131
genre=Romance	50
genre=Children's	89
genre=Documentary	123
genre=Sci-Fi	46
genre=Horror	262
genre=Western	33
genre=Mystery	36

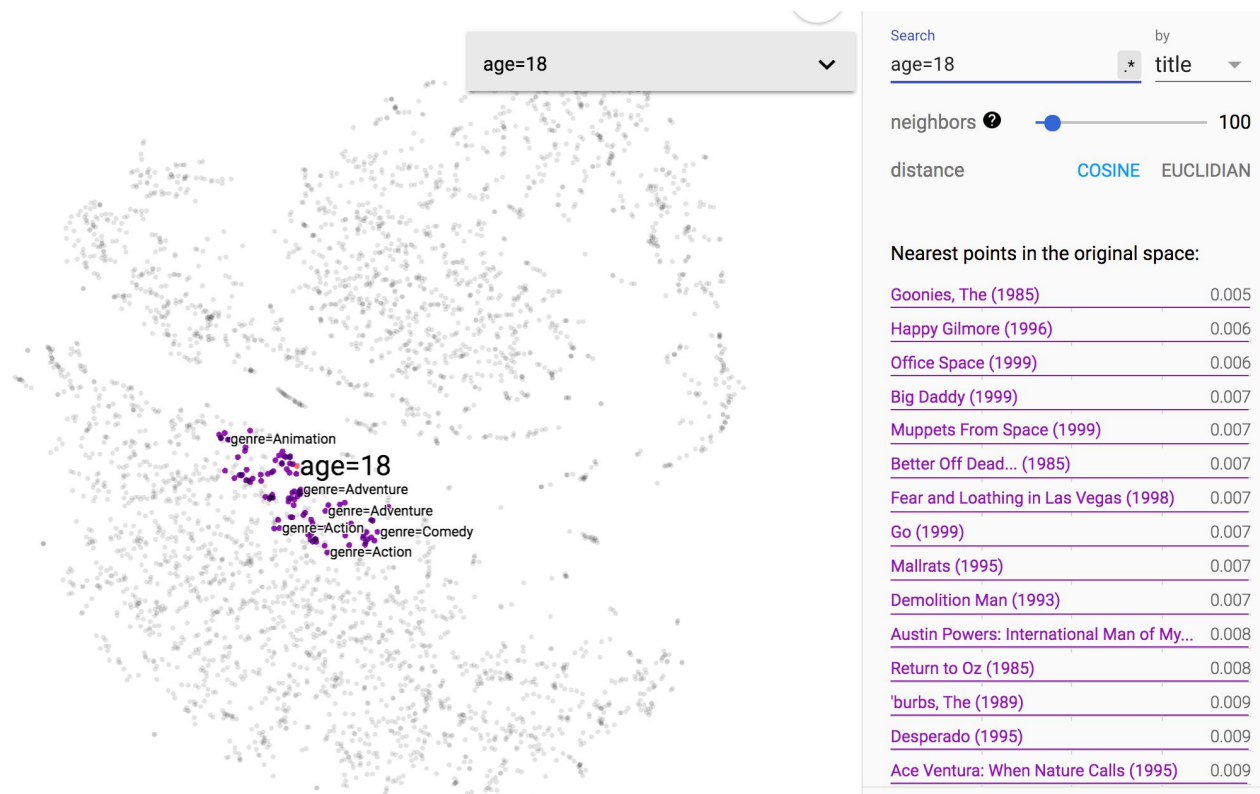
Recommendations as a nearest neighbour search



Recommendations as a nearest neighbour search



Recommendations as a nearest neighbour search



Conclusions and openings

Python ecosystem with TF strikes the right balance

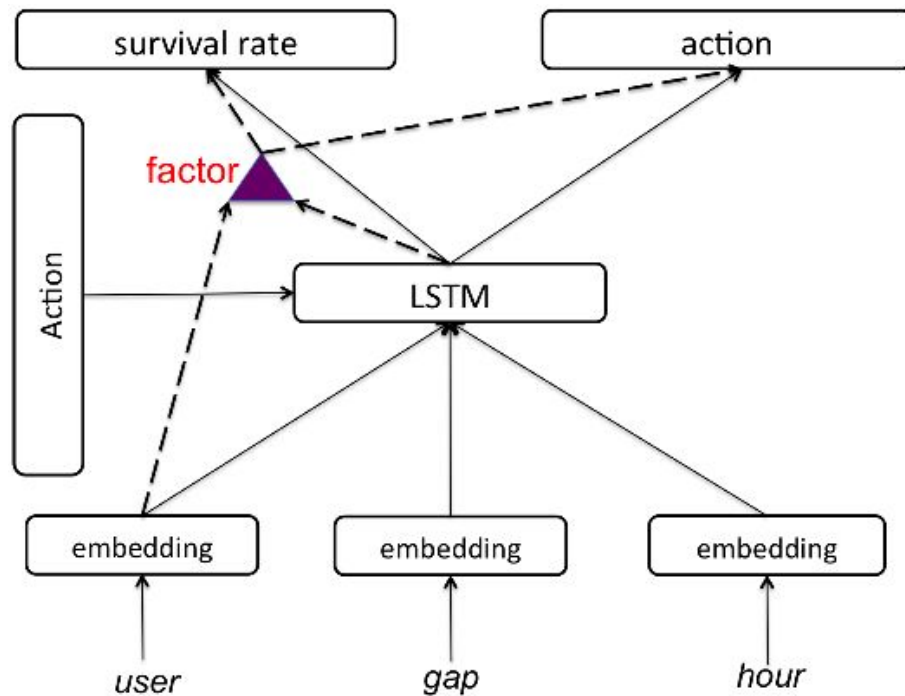
- Possible to implement state of the art of recommender systems
- With guarantees wrt scaling & deployment

Lots more to cover!

- Learning to rank: [Neural net approach with a pairwise \(BPR\) ranking loss](#)
- Wide and deep recommenders, [Factorisation machines with Tensorflow...](#)
- Sequence modelling and time-dependencies

Neural Survival Recommender (2017)

- RNNs to model sequential actions
- Jointly predicting **When** and **What** will be the user's next action



More bibliography

- Fast single core ALS <http://www.benfrederickson.com/matrix-factorization/>
- [Large scale ALS harnessing GPUs](#)
- [AutoRec: Autoencoders Meet Collaborative Filtering](#)
- [Collaborative Deep Learning for Recommender Systems](#)
- [Generating Recommendations at Amazon Scale with Apache Spark and Amazon DSSTNE](#)

Links to my notebooks:

- [ML1m-history2multiclass](#)
- [Movielens-binary](#) and [movielens-fm](#) using [lightfm](#)

Collaborative Denoising Auto-Encoders

